

Linear Classification 1 - LDA

Julian Karch




Universiteit
Leiden
The Netherlands

Overview

- Intro Generative Classification + Naïve Bayes
- Linear Discriminant Analysis
- Quadratic Discriminant Analysis and Regularized LDA

Intro + Naïve Bayes




Universiteit
Leiden
The Netherlands


Classification

- Last week regression: $Y \in \mathbb{R}, \hat{Y} = f(X)$
- This week classification: $Y \in \{1, \dots, K\}, \hat{Y} = \mathcal{C}(X)$

K Nearest Neighbors



K Nearest Neighbors



Expected Prediction Error Regression

- Last week: For regression, we want f that minimizes
$$MSE_{pred} = \mathbb{E}[(f(X) - Y)^2]$$
- More general: $EPE = \mathbb{E}[L(f(X), Y)]$
- L is a **loss function** that quantifies how “bad” wrong predictions are
- We get MSE_{pred} by using $L_{sq}(\hat{Y}, Y) = (\hat{Y} - Y)^2$
- Another option: $L_{abs}(\hat{Y}, Y) = |\hat{Y} - Y|$

Expected Prediction Error Classification

$$\text{EPE} = \mathbb{E}[L(C(X), Y)]$$

- Most popular loss is 0-1 loss

$L_{01}(C(X), y) = 0$ if $y = C(x)$, otherwise $L_{01}(C(X), y) = 1$

Posterior Probability


- $P(\text{no rain}|\text{data})=80\%$
- More generally: $P(Y|X)$ is called *posterior probability*

☐ Would you act like it would rain tomorrow?

Bayes Optimal Classification

- We minimize EPE by assigning the class that has the highest posterior probability
- This is called the *Bayes classifier* and is optimal in the sense that it minimizes EPE (for 0-1 loss) among **all** classifiers
- Note that sometimes 0-1 loss is not what we want
 - Example: falsely predicting “no rain” can be more costly than falsely prediction “rain”

From Data to Posterior



Universiteit
Leiden
The Netherlands

Corona Test

- Corona Test X (positive or negative), Health Status Y (corona or healthy)
- $P(X = \text{positive} | Y = \text{Corona}) = 0.9$
- $P(Y = \text{Corona}) = 0.02, P(Y = \text{Healthy}) = 0.98$
- $P(X = \text{positive} | Y = \text{Healthy}) = 0.05$

Q

- **What is $P(X = \text{positive}, Y = \text{Corona})$?**
- **What is $P(X = \text{positive}, Y = \text{Healthy})$?**
- **What would you classify (assume 0-1 loss)?**

Bayes Rule

- Want: $P(Y = \text{Corona} \mid X = \text{positive})$
- $P(Y = \text{Corona} \mid X = \text{positive}) = \frac{P(X=\text{positive}, Y=\text{Corona})}{P(X=\text{positive})}$
- $P(X = \text{positive})$ does not depend on Y
- \rightarrow Assigning to class y for which $P(X = \text{positive}, Y = y)$ is highest is equivalent to assigning for which $P(Y = y \mid X = \text{positive})$ is highest

Bayes Rule

- $P(X = \text{positive}) = P(X = \text{positive}, Y = \text{Corona}) + P(X = \text{positive}, Y = \text{Healthy})$
- $P(Y = \text{Healthy} | X = \text{positive}) = \frac{.049}{.018 + .049} = 0.73$

Intermediate Summary

- Posterior probabilities $P(Y|X)$ enable easy classification
- Bayes rule allows obtaining posterior probabilities from prior probabilities $P(Y)$ and likelihoods $P(X|Y)$


$$P(Y|X) = \frac{P(X|Y)P(Y)}{P(X)}$$

- Note that this optimality requires that we know the likelihood ($P(X|Y)$) and prior probabilities [$P(Y)$]
- In practice, we especially virtually never know the likelihood $P(X|Y)$


Estimating Prior and Likelihood

- Example: Classify male and female based on their heights (have training set)
- Use relative frequencies for estimating prior probabilities $P(Y = Female)$ and $P(Y = Male)$ (assume $P(Y = Female) = 60\%$)
- For likelihood, make assumption: Within groups heights are normally distributed
- Via standard estimation: $P(X|Y = Male) = N(183, \sigma^2 = 49)$, $P(X|Y = Female) = N(176, 81)$

Likelihood (Prob. of height given group)



Posterior (Prob. of group given height)



Multiple Features: Naïve Bayes

- If we have multiple features, we need a model for the joint distribution
- The easiest / naïve approach is to assume that all features are independent
- Formally, naïve Bayes assumes:
$$P(X = x|Y = k) = \prod_{i=1}^p P(X_j = x_j|Y = k)$$
- For each feature, a model is required
- Typically those models are parametric distribution families (Gaussian, uniform etc.)
- For those families, parameters are estimated per class (class-specific mean and variance, for example)

Linear Discriminant Analysis




Universiteit
Leiden
The Netherlands


LDA $p=1$

- The same as the height example, but assumes equal variances across classes

LDA Likelihood



LDA Posterior

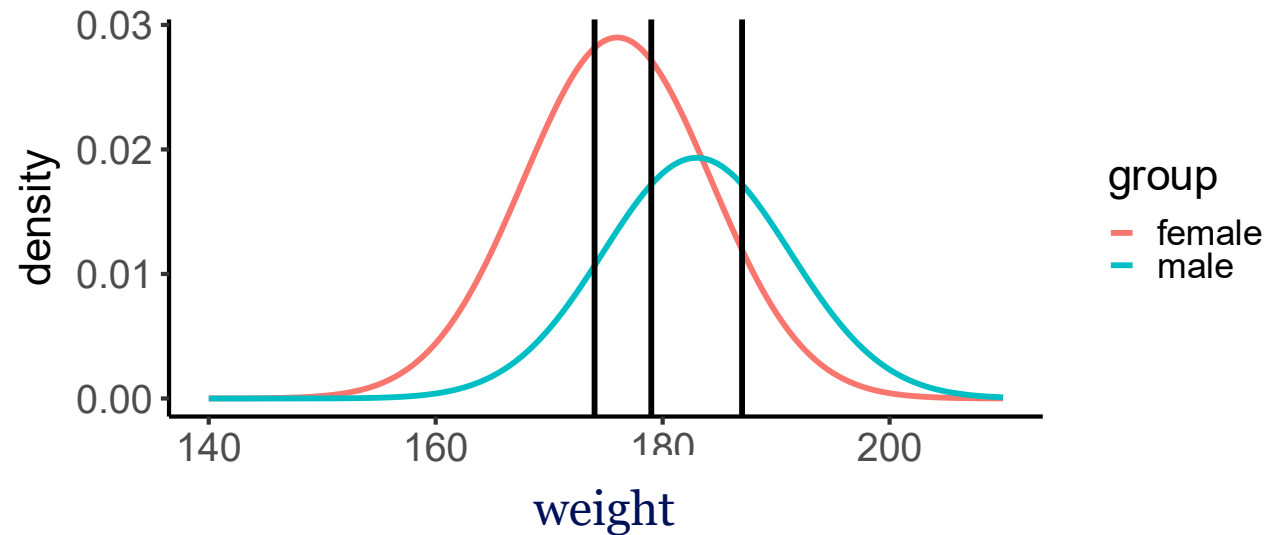



LDA $p > 1$

- Idea: Fit a univariate LDA per dimension

weight

LDA $p > 1$

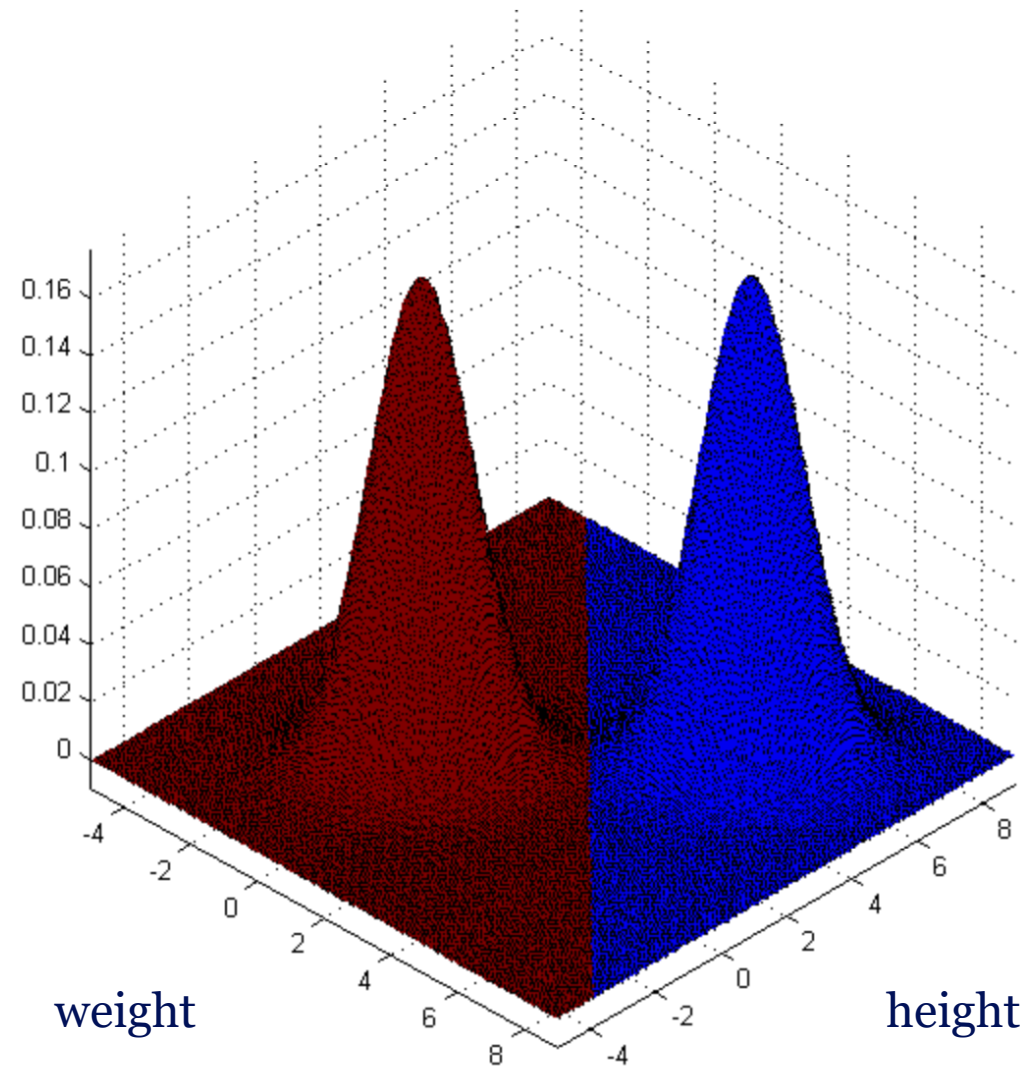


LDA $p > 1$

- Assumes independence of features: often heavily violated
- Fit **multivariate** Gaussian to each class
- **Assume equal covariances**

weight

LDA $p > 1$



Linear Discriminant Analysis

Q


Which color does LDA classify this point to?

1

RED

2

Blue



Linear Classifier

- x is feature vector with $x = [x_1, x_2, \dots, x_p]^\top$
- A classifier is called linear if its decision function is of the form


$$C(x) = \text{sign} \left(\beta_0 + \sum_{j=1}^p \beta_j x_j \right)$$

- Defining $\tilde{x} = [1, x]^\top$ this can more compactly be expressed as

$$C(x) = \text{sign} (\beta^\top \tilde{x})$$


- It follows directly that the decision boundary is linear (linear / hyper-plane)

Intuition LDA Linear



Quadratic Discriminant Analysis

- Drop the assumption of equal covariance matrices
- Leads to quadratic decision boundaries
- Decreases bias at the cost of increased variance



To focus in book (for LDA)

- Formulas for LDA and QDA
- Discriminant function

Bias and Variance

- LDA generally low variance, high bias
- If within-class distribution normal and same covariance \rightarrow no bias
- QDA higher variance but lower bias

Performance Metrics



Universiteit
Leiden
The Netherlands

Problems with Accuracy

- $\text{Accuracy} = \frac{\text{\#classifier correct}}{\text{\#total predictions}}$
- Want to predict patient has Corona yes/no (positive / negative)
- Assume: at any given time 2% of the tested people are positive
- $C(x) = \text{"negative"}$ has accuracy 98%
- For unbalanced classes, accuracy tends to be misleading

Confusion Matrix

- TP: Number of samples that are *correctly* classified as positive
- TN: Number of samples that are *correctly* classified as negative
- FP: Number of samples that are *incorrectly* classified as positive
- FN: Number of samples that are *incorrectly* classified as negative

	Predicted Positive	Predicted Negative
Actually Positive	True Positive (TP)	False Negative (FN)
Actually Negative	False Positive (FP)	True Negative (TN)

Hours of Fun With 4 Numbers

- Accuracy: $\frac{\# \text{classifier correct}}{\# \text{total predictions}} = \frac{TP+TN}{TP+FP+TN+FN}$

“The probability of correctly classifying a random person”

- Sensitivity: accuracy within the positive class: $\frac{TP}{TP+FN}$

“The probability of the test to be positive if the patient has Corona”

- Specificity: accuracy within the negative class: $\frac{TN}{TN+FP}$

- “The probability of the test to be negative if the patient is healthy”

Hours of Fun With 4 Numbers

- Positive Predictive Value: accuracy within the positive predictions

$$\frac{TP}{TP+FP}$$

“Probability of having Corona after having been tested positive”

- Negative Predictive Value: accuracy within the negative predictions:

$$\frac{TN}{TN+FN}$$

“Probability of being healthy after having been tested negative”

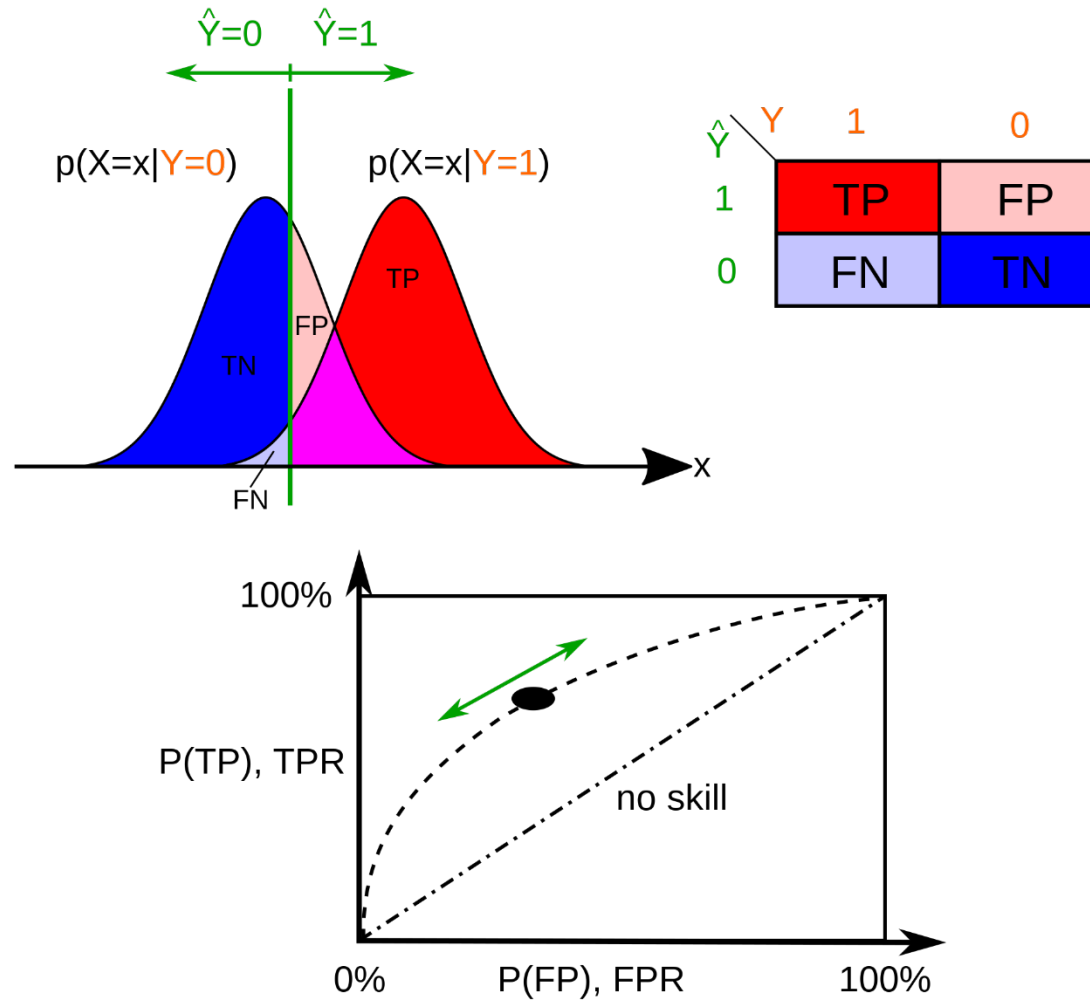
Hours of Fun With 4 Numbers Continued

- True positive rate (TPR) = sensitivity = Recall, True negative rate (TNR) = specificity
- Balanced accuracy = $\frac{TPR+TNR}{2} = \frac{sensitivity + specificity}{2}$ = average of within class accuracies
- Positive Predictive Value = Precision
- F-Measure = harmonic mean of precision and recall $\frac{2}{\frac{1}{precision} + \frac{1}{recall}}$
- Harmonic mean is closer to the smaller value
 - $H(0.5,0.5) = 0.50$
 - $H(0.5,0.6) = 0.54$
 - $H(0.5,1) = 0.67$

TP vs FP Tradeoff

- True positive rate (TPR)= accuracy within the positive class: $\frac{TP}{TP+FN}$
 - “The probability of the test to be positive if the patient has Corona”
- True negative rate (TNR)=accuracy within the negative class: $\frac{TN}{TN+FP}$
 - “The probability of the test to be positive if the patient has Corona”
- False positive rate (FPR)=inaccuracy within the negative class:
 $1 - TNR$

ROC Curve



ROC Curve

- Shows TPR and FPR for every threshold
- Optimally $TPR = 1$ and $FPR = 0$
 - Corresponds to horizontal line at 1
- Random guessing \rightarrow diagonal
- Above the diagonal \rightarrow better than random guessing
- Below the diagonal \rightarrow worse than random guess
- AUC

When to Use What

- Have to distinguish: Evaluation vs. Optimization
- For evaluation, I recommend using many different methods, as quantify many different, important aspects of performance
- For selection: Select based on what is most important

Performance Measures for Selection

- If every misclassification is equally “bad” and the distribution of classes is roughly equal or the same between training and test → Accuracy
 - Example: Is a car red or orange?
- Same conditions but unequal class sizes and likely difference between training and test → Balanced Accuracy
- If misclassifying the positive class is much worse than the negative class → TPR
 - Example: Corona test

Performance Measures for Selection

- If misclassifying the negative class is worse than the negative class \rightarrow FPR
 - Example: Spam detection
- If misclassifying the positive class is worse, you also want to consider the negative, and unbalanced class \rightarrow F1-score
 - AUC: When you do not care about a particular threshold but the performance at many thresholds
 - Example: Evaluation of probabilistic classifier

Question Performance Metrics

Q Consider a classifier that always predicts the positive class. What is its TPR and what is its FPR rate?

Hours of Fun with 4 Numbers (the end)

- Many more measures
 - https://en.wikipedia.org/wiki/Evaluation_of_binary_classifiers